

Title of your thesis

L^AT_EX

Piotr Janisz

Faculty of Electrical Engineering, Automatics, Computer Science and
Electronics
University

A thesis submitted for the degree of

Philosophiæ Doctor (PhD), DPhil,..

year month

Abstract

Computing power growth allows more complex and accurate simulation techniques to be implemented in real-time applications. An example of those techniques are particle-based methods for simulation of fluids. They can provide a new level of realism into computer games. Nowadays most fluids in games are simulated using height field. Although using this method realistic waterbodies (like oceans or ponds) can be simulated, it's hard to achieve effects such as splashing or flooding. Those can be easily simulated with particle-based methods such as Smoothed Particle Hydrodynamics (SPH).

In this paper I would like to present realistic model of fluid that can be used in real time application, especially in computer games. Emphasis will be placed on realistic rendering output from particle simulation. For water simulation NVIDIA PhysX SDK will be used. For rendering particles I will use two different approaches: screen space rendering and isosurface extraction. Second algorithm is described in [reference] and is running on cpu. Authors presented performance of this algorithm running on one CPU core and I will present multithread implementation and it's performance analysis. At the end of this paper I will also present comparison of those two rendering techniques.

Contents

1	Introduction	1
2	Rendering techniques	3
2.1	Screen space	3
2.1.1	Surface depth	4
2.1.2	Smoothing	5
2.1.2.1	Gaussian smoothing	5
2.1.2.2	Curvature flow smoothing	8
2.1.3	Thickness	9
2.1.4	Rendering	9
2.2	Isosurface extraction	9
2.2.1	Overview	9
2.2.2	Block subdivision	10
2.2.3	Block processing	11
2.2.4	Lookup cache	12
2.2.5	Multithreading	12
	Bibliography	13

CONTENTS

1

Introduction

1. INTRODUCTION

2

Rendering techniques

Output from particle base simulations is list of particles containing positions. It can also contain additional parameters - for example PhysX fluid simulation returns velocity, lifetime and density in addition to position for each particle. Traditionally triangle meshes are used for rendering objects. Using this approach requires constructing fluid triangle mesh for given particle positions. This method is called isosurface extraction and will be described in second section of this chapter. Another possibility is to render each particle as a point (or a billboard in general) with opacity so that when large amount of particles is rendered the result would give illusion of a smooth surface. Such technique has one major drawback - it does not produces surface prevents us from adding effects of reflection and refraction. Similar technique is to extract visible surface by rendering each particle as a sphere into depth buffer. This will be described in first section.

2.1 Screen space

As mentioned before this approach extracts visible surface by rendering each particle into depth buffer. High level overview of this method is presented on figure 2.1 and consists of following steps (2): rendering depth texture (section 2.1.1) and thickness textures (section 2.1.3), depth smoothing (section 2.1.2) and assembling fluid surface with rest of the scene (section 2.1.4).

2. RENDERING TECHNIQUES




Figure 2.1: Screen Space Fluid Rendering - The figure shows high level overview of the method, taken from (1)

2.1.1 Surface depth

To obtain fluid surface visible from the viewpoint of camera each particle is rendered as a sphere into depth texture [rysunek]. At each pixel only closest value is kept using hardware depth test. To avoid rendering large amount of geometry each particle is rendered as a point sprite and it's depth is generated in fragment shader. This common technique speeds up rendering process significantly as well as improves quality of rendered spheres (as can be seen on figure 2.2) because depth values are generated for each pixel. Rendering spheres as point sprites is 10 to 100 times faster comparing to rendering them as triangle meshes (see table 2.1).

Table 2.1: Comparison of sphere rendering methods

Method	Frames per second
Mesh, 128 triangles	20
Mesh, 512 triangles	6
Mesh, 2048 triangles	1
Point Sprites	200




Figure 2.2: Spheres rendered with different methods - 1 - mesh with 128 triangles, 2 - mesh with 512 triangles, 3 - mesh with 2048 triangles, 4 - point sprite with normals generated in fragment shader

2.1.2 Smoothing

Although previous step produces surface it's quality is not sufficient. As can be seen on figure 2.3 individual spheres can be seen giving fluid surface gelly-like appearance. To remove this artifact some kind of smoothing has to be applied. Section 2.1.2.1 will describe gaussian smoothing and section 2.1.2.2 curvature flow smoothing.

2.1.2.1 Gaussian smoothing

Most obvious way to smooth values in depth texture is to apply gaussian filter. It's easy to implement and can be computed fast due to it's linear separability. However this filter produces undesired effect of blending drops of fluid with background surfaces (see figure 2.4). Thus edge-preserving filters needs to be used which are also called bilateral filters. Bilateral Gaussian filter is a modification that changes wages of pixels depending on difference between their tonal value ($I(s)$) and tonal value of central pixel ($I(s_0)$). This can be described by following formula (from (3)):

$$O(s_0) = \frac{\sum_{s \in S} f(s, s_0) I(s)}{\sum_{s \in S} f(s, s_0)} \quad (2.1)$$

2. RENDERING TECHNIQUES




Figure 2.3: Fluid surface rendered without smoothing phase - Individual spheres can be seen, giving surface unnatural appearance

TODO

Figure 2.4: Fluid surface rendered with Gaussian smoothing - From left: depth image, depth image after applying Gaussian filter, diffuse shaded surface

where

$$f(s, s_0) = g_s(s - s_0)g_t(I(s) - I(s_0)) \quad (2.2)$$

is the bilateral filter for the neighborhood around s_0 . g_s is a spatial weight, g_t is a tonal weight and they both are Gaussian functions:

$$g_s(s) = g(x, \sigma_s)g(y, \sigma_s) \quad g_t(I) = g(I, \sigma_t) \quad (2.3)$$

As can be seen on equation 2.2 only change in bilateral filter (in comparison to regular bilateral filter) is introduction of tonal weight g_t . This change makes filter space-variant - that means it can't be computed as a product of two one dimensional filters (g_s in equation 2.3). Computational complexity of space-variant filters is $O(Nm^d)$ comparing to only $O(Nmd)$ for space-invariant (d is image dimensionality, m is the size of filtering kernel and N is the number of pixels in the image). For smoothing fluid surface kernel with $m = 20$ must be used, which gives for 2 dimensional image about 400 operations for each pixel when using bilateral filtering. In comparison separable implementation requires 40 operations. It turns out however that computing bilateral filter as it was a space-invariant gives good approximation for real time applications (see figure 2.5). Approximation gives some artifacts but they are not visible when fluid is moving and other effects (reflection and refraction) are applied.

TODO

Figure 2.5: Comparison of normal bilateral filter with it's separable approximation - Upper row presents images for bilateral filter and bottom row presents separable approximation.

Bilateral filter has some undesired effect when applied on z-buffer output. The problem is that depth values are not evenly distributed. This mean that difference

2. RENDERING TECHNIQUES

between two fragments depth values is dependent on their distance to the viewer. As can be seen on equation 2.2 bilateral weights takes into account difference between depth values. In result the further surface lies from camera the less edges are preserved (see figure 2.6). Alternative technique to Z-buffering is W-buffering which offers better distribution of depth values (4). Some hardware supports W-buffering but most of it doesn't and it doesn't seems to be supported in future. A workaround to this problem was presented in (5). It customizes projection transformation in vertex shader in way that depth buffer values have linear distribution at the end.

TODO

Figure 2.6: Bilateral filtering results for different distances to camera - The more distant fluid is from viewer the more smoothed its surface is. Also less edges are preserved due to z-buffer used resulting in more particles got blended into background surfaces

Another issue with smoothing using fixed kernel size filter is that further surfaces are more smoothed than closer surfaces (see figure 2.6). Solving that problem would require using filter with variable kernel size, depending on fragment distance from viewer. As implementing this efficiently on graphic hardware is hard and the effect is not so irritating this won't be considered any more. Section 2.1.2.2 will describe another smoothing technique that doesn't suffer from this problem.

2.1.2.2 Curvature flow smoothing

This technique was presented in (2).

2.1.3 Thickness

This step is computed to determine how opaque is surface in given point. It is achieved by rendering particles as circles into depth buffer with additive blending enabled. Resulting thickness texture is then smoothed with Gaussian filter (this time regular one). In order to speed up rendering process this texture can be rendered with lower resolution and then applied with linear interpolation.

2.1.4 Rendering

Last step assembles fluid surface with background image. As an input it takes fluid depth texture, thickness texture and texture with rest of the scene rendered. [TODO normal reconstruction, refraction, reflection, thickness]

2.2 Isosurface extraction

Classic algorithm for isosurface extraction is marching cubes (see (6)). It takes 3d array with scalar field values as an input and produces list of triangles. Algorithm divides space into cubes, and proceeds through scalar field taking one cube at a time (each cube consists of 8 vertices with scalar field values). Field value at each cube vertex is tested to be above or under given threshold and then appropriate triangle configuration is taken from lookup table (there are 256 triangle configurations).

This method is not suitable for extraction surface from particles. Firstly scalar field values at cube corners are not given. Secondly visiting each cube is not efficient as most of cubes doesn't contain surface. Several variations of marching cube algorithm were created to overcome those problems. The algorithm used is a multithreaded version of the one presented in (7). This is a surface following version of marching cubes with a fast particle lookup cache technique.

2.2.1 Overview

Algorithm takes as an input a list of particles positions, particle radius of influence (R_c), isosurface threshold and produces list of triangles. Particle radius of influence is something different than particle size (or simply particle radius - R), because its value is usually greater. Every particle produces a field that has non-zero values within R_c

2. RENDERING TECHNIQUES

from particle center. The field value in given point of space (p) is a sum of field values generated by all particles within R_c . This can be represented by:

$$F(p) = \sum_{s \in D_{R_c}(p)} f(dst(s, p)) \quad (2.4)$$

where s is a center of particle, $D_{R_c}(p)$ is a disc with p as a center and radius R_c . $f(r)$ is a field function of single particle. It can be any function that is radially monotonic, continuous and has non-zero values within R_c from particle center. Following function is often used as it can be computed efficiently with a few operations:

$$f(r) = \begin{cases} (\frac{r}{\sqrt{2}R_c})^4 - (\frac{r}{\sqrt{2}R_c})^2 + 0.25 & \text{if } r < R_c \\ 0 & \text{otherwise} \end{cases} \quad (2.5)$$

When $R_c > R$ neighboring particles combines into one smooth surface (see 2.7). This is a common technique known as metaballs (see TODO cytat). The biggest problem here is to efficiently compute field values at corners and to traverse space in most optimal way. Those topics will be described in following subsections.




Figure 2.7: Illustration of combining particles - SOMETHING

2.2.2 Block subdivision

Space is divided into blocks containing cubes. Such a division aims for:

- Reduce memory consumption - smaller blocks contains less particles and cubes.
- Allowing for parallelize algorithm in an easy way - because processing one block is an independent of processing other blocks.

Blocks are expanded to contain also particles that lies outside them but may influence the field values inside block. That means the block must also contain particles lying within R_c distance from it - this area is called margin and has special treatment, which will be described in next section. Expanding blocks makes parallel processing very easy as there are no dependencies between them.

TODO imgae of block

Block subdivision is a first step of the algorithm. For every particle containing blocks are computed. One particle can be contained in at most 8 blocks due to blocks overlapping. Each block has a list of particles lying within its boundary.

2.2.3 Block processing

Processing of block starts from building a particle lookup cache which will be described in section 2.2.4. Next surface following marching cube algorithm is run inside block. The algorithm is named marching slices in (7). Block containing $n \times n \times n$ cubes is divided to n slabs and $n+1$ slices (TODO rysunek). Division is made in parallel to XY plane. Each slab contains $n \times n \times 1$ cubes and each slice contains $(n+1) \times (n+1) \times 1$ cube corners. Slices lies between slabs and serves as a cache for field values computed at corners. In addition neighboring slabs share slices that lies between them. Each slab has a list of cubes lying within it that should be visited - a todo list.

To speed up a marching cubes algorithm we would like to traverse only cubes that contains surface. This requires finding so called seeding cubes - initial cubes at which algorithm will start traversing block. We then choose one slab tat contains one or more seed cubes and pick up one cube. Field values at cube corners are calculated and put into slice caches, surface is polygonized and neighboring cubes that contain surface are aded to todo lists of [odpowiadajacy] slabs.

If we pick up a random slabs and cubes each time then we have to keep all cached values in slices while processing a cube. The idea of marching slices is to decrease memory used by cache by traversing cubes slice by slice and deallocate caches as soon as possible. This can be done by always picking the lowest slab with not empty todo list and finding seeding cubes in the lowest possible slabs.

The optimum would be to find surfaces all local minimas - then we would have to store only cache in slices below and above current slab. However this is approach would

2. RENDERING TECHNIQUES

be time consuming. (7) used heuristic approach with complexity $O(P)$ where P is a number of particles. The algorithm is as follows:

```
for every particle:
    take a corner closest to particle center
    while  $F(c) > \text{threshold}$  and  $|c - c_{\text{start}}| < 2$  do
         $c = \text{corner below } c$ 
    if  $F(c) < \text{threshold}$ 
        add cube to it's slab todo list
```

This method will not find local minimas every time as it always goes down from the centre of a particle. Figure 2.8 shows the case when algorithm (TODO ref) fails to find a local minimum.

TODO

Figure 2.8: Finding local minima - In this example algorithm fails to find local minimum

2.2.4 Lookup cache

To compute field value for a given corner all particles that lies within r_c from it have to be found. This task can be accomplished by building space partitioning data structure like octtree (TODO moe odnoniki). However (7) presented another to accomplish this task. Presented method does use only integer comparison

2.2.5 Multithreading

diagram

Bibliography

- [1] SIMON GREEN. **Screen Space Fluid Rendering for Games**. Game Developers Conference, 2010. 4
- [2] MIGUEL SAINZ WLADIMIR J. VAN DER LAAN, SIMON GREEN. **Screen space fluid rendering with curvature flow**. *Proceedings of the 2009 symposium on Interactive 3D graphics and games*, 2009. 3, 8
- [3] LUCAS J. VAN VLIET TUAN Q. PHAM. **Separable Bilateral Filtering for Fast Video Preprocessing**. *Multimedia and Expo, 2005. ICME 2005. IEEE International Conference on*, 2005. 5
- [4] JASON GREGORY. **Occlusion and the Depth Buffer**. In *Game Engine Architecture*, pages 442–444. A K Peters, Ltd., 2009. 8
- [5] ROBERT DUNLOP. **Linearized Depth using Vertex Shaders**. Article on X-Zone service, 2006. 8
- [6] HARVEY E. CLINE WILLIAM E. LORENSEN. **Marching cubes: A high resolution 3D surface construction algorithm**. *Proceedings of the 14th annual conference on Computer graphics and interactive techniques*, Vol. 21, No. 4., 1987. 9
- [7] KEN BIRDWELL ILYA D. ROSENBERG. **Real-time particle isosurface extraction**. *Proceedings of the 2008 symposium on Interactive 3D graphics and games*, 2008. 9, 11, 12